

This member-only story is on us. [Upgrade](#) to access all of Medium.

Member-only story

Boost Your Django App's Speed: The Ultimate Guide to Asynchronous Queries

QuerySet and Manager Methods in Asynchronous Code



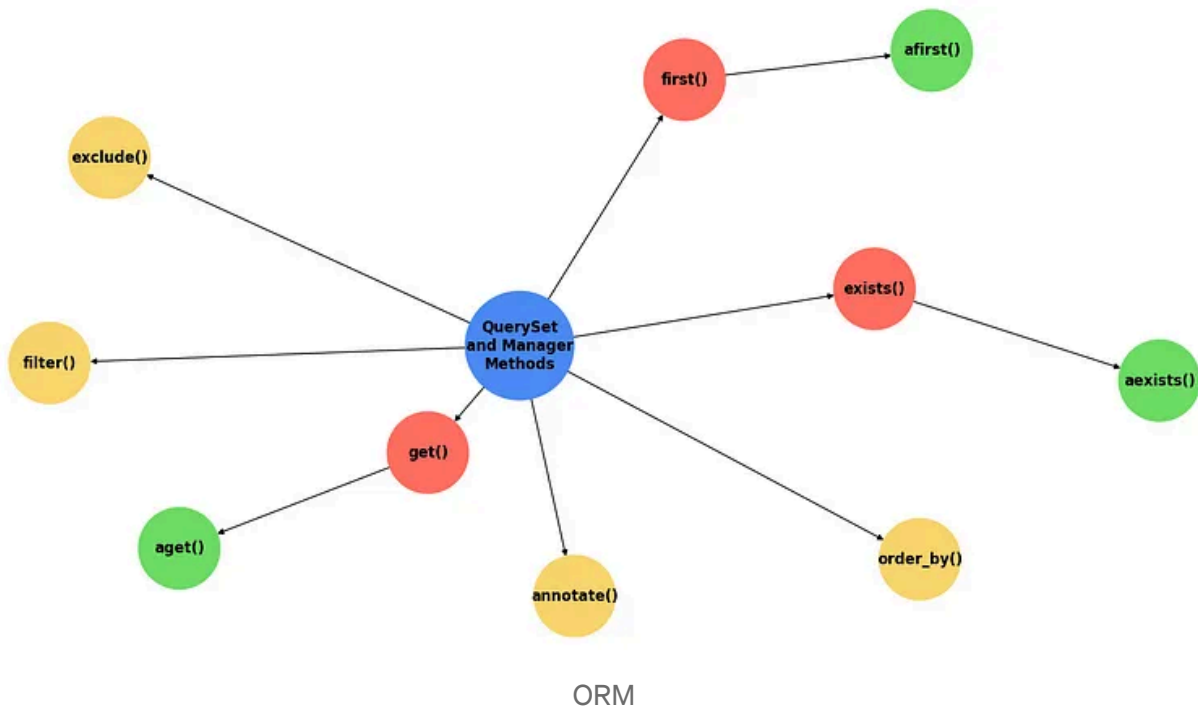
Ewho Ruth · Follow

9 min read · Oct 17, 2024

Listen

Share

More



Introduction

When working with Django's ORM in asynchronous views or code, it's crucial to distinguish between methods that **force execution** of the `QuerySet` and those that do

not. This will help you determine which methods require the asynchronous variants.

Blocking vs. Non-blocking Methods

1. Blocking Methods:

- Methods like `get()` and `first()` execute the query immediately and return a single result or raise an exception. These methods are blocking and should be replaced with their asynchronous counterparts in async code.
- Asynchronous equivalents include `aget()` and `afirst()`.

2. Non-blocking Methods:

- Methods like `filter()` and `exclude()` do not execute the query right away. Instead, they return a new `QuerySet` that can be further refined without making immediate database calls. These methods can be safely used in asynchronous code.

Blocking Methods

These methods execute the query immediately and retrieve data from the database synchronously. When working in asynchronous code, these must be replaced by their asynchronous equivalents to prevent blocking. Examples include:

- `get()` → Asynchronous equivalent: `aget()`
- `first()` → Asynchronous equivalent: `afirst()`
- `exists()` → Asynchronous equivalent: `aexists()`

Non-blocking Methods

These methods do not hit the database immediately. They return new `QuerySet` objects that can be further refined without making database calls. These are safe for use in both synchronous and asynchronous code:

- `filter()`
- `exclude()`
- `annotate()`
- `order_by()`

Use Cases for Async Methods

Asynchronous methods can be particularly useful in scenarios where your application must handle a large number of concurrent I/O-bound tasks.

Here's an example of where they shine:

Handling High-Concurrency API Requests

Consider an API that needs to retrieve multiple user profiles from the database. If each query takes 200ms and hundreds of requests come in simultaneously, a synchronous implementation would struggle to handle the load.

Asynchronous methods allow the server to handle many concurrent requests without blocking.

```
async def get_user_profiles(request):
    usernames = request.GET.getlist('usernames')
    # Fetch multiple user profiles concurrently
    users = await User.objects.filter(username__in=usernames).all()
    return JsonResponse([user.serialize() for user in users], safe=False)
```

By using `await`, your code can efficiently retrieve multiple users concurrently, improving throughput and response time.

Performance Implications of Asynchronous Methods

I/O-bound Tasks

Asynchronous methods excel when dealing with I/O-bound tasks like database queries or external API calls. They allow the application to release the CPU while waiting for the I/O operation, so the event loop can continue handling other tasks.

This can lead to substantial performance gains, especially under high concurrency.

However, it's important to recognize that async code isn't always beneficial. In **CPU-bound tasks**, which require significant processing power, the Global Interpreter Lock (GIL) in Python still limits the execution of threads, making async code less effective.

Performance Gains in I/O-bound Operations

For operations like querying databases or making external API requests, using async can drastically reduce wait times, improving response times in high-traffic environments.

Best Practice: When to Use Asynchronous Methods

1. Use Async for I/O-bound Operations

Async shines in I/O-bound tasks such as database queries or calling external APIs. It allows the server to handle other requests while waiting for a database response, improving the overall throughput. Use asynchronous methods for actions like:

- Fetching records from the database
- Making external HTTP requests
- Handling file I/O operations (e.g., uploads)

```
async def get_user_profiles(request):  
    usernames = request.GET.getlist('usernames')  
    # Fetch profiles concurrently
```

```
users = await User.objects.filter(username__in=usernames).all()
return JsonResponse([user.serialize() for user in users], safe=False)
```

This approach improves response times, especially when your application deals with many concurrent requests.

2. Avoid Using Async for CPU-bound Operations

For CPU-intensive tasks such as data processing or calculations, async won't offer performance benefits due to Python's Global Interpreter Lock (GIL). For these, consider using multi-threading or offloading the work to a task queue (e.g., Celery).

Alternatives to Asynchronous Queries

1. Query Caching

If you expect certain database queries to run repeatedly, consider caching the results using Django's cache framework instead of fetching data asynchronously every time. Caching can improve performance by reducing the number of database hits.

```
from django.core.cache import cache

async def get_cached_user_profile(user_id):
    user = cache.get(f'user_{user_id}')
    if not user:
        user = await User.objects.aget(id=user_id)
        cache.set(f'user_{user_id}', user, timeout=60*60) # Cache for 1 hour
    return user
```

This method trades off the complexity of async handling with performance gains from caching, reducing database load.

2. Task Queues for Long-running Operations

For operations that might take time (such as sending emails or processing large datasets), consider using a task queue like Celery instead of async views. This keeps your application responsive while offloading long-running tasks.

Open in app ↗

Medium



Search



B

```
return JsonResponse({"message": "User created"})
```

This approach ensures that heavy operations don't slow down your primary web server.

When Not to Use Async

For smaller applications with limited I/O or concurrency needs, introducing async code may not provide significant performance benefits. In such cases, the added complexity of asynchronous programming can outweigh any potential improvements.

Addressing Potential Challenges and Limitations

- Handling Exceptions in Asynchronous Code

When working with async methods, you may encounter exceptions such as `TimeoutError` or other database-related issues. Here's how to manage these exceptions effectively:

```
try:
    user = await User.objects.aget(id=user_id)
except User.DoesNotExist:
    return JsonResponse({"error": "User not found"}, status=404)
except asyncio.TimeoutError:
    return JsonResponse({"error": "Request timed out"}, status=504)
```

In this example, a `TimeoutError` is handled gracefully, ensuring the application remains responsive even if a query takes too long.

- **Managing Timeouts and Retries**

Timeouts are a common concern when working with async code. Setting a timeout for database operations can prevent the application from waiting indefinitely:

```
import asyncio

async def get_user_data(user_id):
    try:
        # Set a timeout for the asynchronous DB query
        return await asyncio.wait_for(User.objects.aget(id=user_id), timeout=5.)
    except asyncio.TimeoutError:
        # Handle the timeout
        return None
```

In this example, a 5-second timeout is set to prevent the query from running indefinitely, and the timeout is handled gracefully.

- **Retry Logic**

In some cases, implementing retry logic with async methods can improve the reliability of your code, especially when dealing with transient failures like database connection issues.

```
import asyncio

async def get_user_data_with_retry(user_id, retries=3):
    for _ in range(retries):
        try:
            return await User.objects.aget(id=user_id)
        except (asyncio.TimeoutError, DatabaseError):
            await asyncio.sleep(1) # Wait before retrying
    return None
```

This approach ensures your application attempts multiple times before giving up, improving its robustness in production environments.

Best Practice: Combine Async and Caching for Optimal Performance

For the best results in performance-critical applications, consider combining asynchronous code with caching. Use async methods to fetch data concurrently, and cache frequently accessed results to reduce the need for future database queries.

```
from django.core.cache import cache

async def get_user_profile(user_id):
    user = cache.get(f'user_{user_id}')
    if not user:
        user = await User.objects.aget(id=user_id)
        cache.set(f'user_{user_id}', user, timeout=60*60) # Cache for an hour
    return user
```


This strategy ensures that your application handles concurrent requests efficiently while reducing database load through caching.

Conclusion: Best Practices Recap

- **Use async selectively:** Focus on I/O-bound operations like database queries and external API calls.
- **Consider caching:** When async queries are repeated, leverage Django's cache framework to improve performance.
- **Avoid async for CPU-bound tasks:** For computationally expensive operations, consider task queues or threading instead of async.
- **Handle exceptions and retries:** Implement proper exception handling, including retry logic for transient errors and timeouts to avoid blocking.
- **Combine async with caching:** For optimal performance, use async methods to fetch data concurrently and cache the results for repeated queries.

Common Errors in Asynchronous Code

One common mistake when using async methods is forgetting to include the `await` keyword. This leads to errors like:

- **“coroutine object has no attribute x”:** This error indicates that the coroutine was not awaited.
- **Coroutine object strings (<coroutine object at ...>):** This happens when the coroutine is not executed because it wasn't awaited.

Always ensure to `await` asynchronous methods to avoid these errors.

Summary

- **Blocking methods** (like `get()` and `first()`) must be replaced by their asynchronous equivalents (`aget()`, `afirst()`) in asynchronous code.
- **Non-blocking methods** (like `filter()` and `exclude()`) are safe to use in async code since they don't trigger database queries immediately.
- Use asynchronous methods in **I/O-bound operations** to improve performance, but understand that for **CPU-bound tasks**, async may not offer substantial gains.
- Handle potential challenges such as **exceptions** and **timeouts** gracefully in async code, and consider implementing **retry logic** for more resilience.
- Always remember to use `await` when calling asynchronous methods to avoid common errors.

How to Identify Method Types

You can refer to the **QuerySet reference** documentation, which categorizes methods into two sections:

- **Methods that return new querysets:** These are non-blocking and safe to use asynchronously (e.g., `filter()`, `exclude()`, `annotate()`).
- **Methods that do not return querysets:** These are blocking and have asynchronous versions (e.g., `get()`, `first()`, `exists()`).

Example of Using Asynchronous Methods

Here's how to handle an asynchronous query correctly:

```
# Asynchronous Query Example
user = await User.objects.filter(username=my_input).afirst()
```

In this example:

- `filter()` returns a `QuerySet`, so you can chain additional methods.

- `afirst()` is used to retrieve the first matching user asynchronously, and you must use `await` to ensure the coroutine is executed properly.

Error Handling with Logging

Handling exceptions properly is essential in async code, but it's equally important to log those errors for debugging and monitoring purposes. Here's an expanded example of how to handle exceptions in asynchronous queries while also logging them for future analysis:

```
import logging
import asyncio

# Set up logging
logger = logging.getLogger(__name__)

async def get_user_with_logging(user_id, retries=3):
    for attempt in range(retries):
        try:
            return await User.objects.aget(id=user_id)
        except User.DoesNotExist:
            logger.error(f"User with ID {user_id} not found.")
            return None
        except (asyncio.TimeoutError, DatabaseError) as e:
            logger.error(f"Error fetching user {user_id} on attempt {attempt+1}")
            if attempt < retries - 1:
                await asyncio.sleep(1) # Retry after delay
    return None
```

Key Points:

- **Logging the error:** The `logger.error` function is used to log detailed information about the exception, including the user ID and attempt number.

- **Re-trying on failure:** By introducing retry logic, the application can attempt to fetch the data multiple times before giving up, logging each attempt.
- **Custom messages:** Logging the specific error (`str(e)`) allows better insight into what went wrong.

This approach ensures that any errors are captured, logged, and available for debugging, making it easier to monitor and trace issues in production systems.

Advanced Asynchronous Features

For those looking to explore more advanced features of asynchronous programming in Python, here are a couple of useful tools you may want to consider:

1. Async Context Managers (`async with`)

Async context managers are used when you want to manage resources that need asynchronous cleanup, such as database connections or network sockets. Here's an example using an async context manager:

```
class AsyncResourceManager:
    async def __aenter__(self):
        print("Resource opened")
        return self

    async def __aexit__(self, exc_type, exc_value, traceback):
        print("Resource closed")

async def handle_resource():
    async with AsyncResourceManager() as resource:
        print("Handling resource")
```

In this example, `async with` ensures that the resource is properly opened and closed asynchronously. This pattern is especially useful when dealing with network connections, file I/O, or any resource that needs proper cleanup.

2. Task Cancellation (`asyncio.ensure_future`)

In some cases, you might need to launch tasks concurrently and have the ability to cancel them. The `asyncio.ensure_future()` function helps manage tasks by scheduling them to run concurrently. You can also cancel them if needed:

```
import asyncio

async def long_running_task():
    try:
        await asyncio.sleep(10)
        print("Task finished")
    except asyncio.CancelledError:
        print("Task was cancelled")

async def main():
    task = asyncio.ensure_future(long_running_task())
    await asyncio.sleep(2)
    task.cancel() # Cancel the task after 2 seconds
    await task # Wait for the task to handle cancellation

asyncio.run(main())
```

Key Points:

- **Task scheduling:** `asyncio.ensure_future()` schedules an asynchronous function to run in the background, allowing you to manage it later.
- **Cancellation:** You can use `.cancel()` to stop a task prematurely and handle the cancellation gracefully with an `asyncio.CancelledError` exception.

By using these advanced async features, you can gain finer control over resource management and task scheduling, further enhancing the flexibility and resilience of your async code.

Method	Blocking	Asynchronous Equivalent
<code>get()</code>	Yes	<code>aget()</code>
<code>first()</code>	Yes	<code>afirst()</code>
<code>last()</code>	Yes	<code>alast()</code>
<code>exists()</code>	Yes	<code>aexists()</code>
<code>count()</code>	Yes	<code>acount()</code>
<code>delete()</code>	Yes	<code>adelete()</code>
<code>update()</code>	Yes	<code>aupdate()</code>
<code>create()</code>	Yes	<code>acreate()</code>
<code>filter()</code>	No	N/A
<code>exclude()</code>	No	N/A
<code>annotate()</code>	No	N/A
<code>order_by()</code>	No	N/A
<code>distinct()</code>	No	N/A

Making queries | Django documentation

The web framework for perfectionists with deadlines.

docs.djangoproject.com

Lambda Function: Python's Playbook: The Pythonista's Guide to Lambda Functions: Write Less, Do More...

Lambda Function: Python's Playbook: The Pythonista's Guide to Lambda Functions: Write Less, Do More eBook : Ewho, Ruth...

www.amazon.in

PER MINUTE: THE POWER OF ONE MINUTE AT A TIME eBook : Ewho, Ruth: Amazon.in: Kindle Store

PER MINUTE: THE POWER OF ONE MINUTE AT A TIME eBook : Ewho, Ruth: Amazon.in: Kindle Store

PER MINUTE: THE POWER OF ONE MINUTE AT A TIME eBook : Ewho, Ruth: Amazon.in: Kindle Store

[Amazon.in: Kindle Store](https://www.amazon.in)

Between Love and Silence eBook : Ewho, Ruth: Amazon.in: Kindle Store

Between Love and Silence eBook : Ewho, Ruth: Amazon.in: Kindle Store

Between Love and Silence eBook : Ewho, Ruth: Amazon.in: Kindle Storeamazon.in

Between Love and Silence Part 2 eBook : Ewho, Ruth: Amazon.in: Kindle Store

Between Love and Silence Part 2 eBook : Ewho, Ruth: Amazon.in: Kindle Store

Between Love and Silence Part 2 eBook : Ewho, Ruth: Amazon.in: Kindle Storeamazon.in

Django

Web3

Web Development

Software Development

Python



Follow

Written by Ewho Ruth

251 Followers · 110 Following

Turning code into fun and sharing my feels!

More from Ewho Ruth



Ewho Ruth



Little Things

Little Things



3d ago



140



6



In Django Unleashed by Ewho Ruth



Understanding Django Field Lookups: A Comprehensive Guide

Introduction

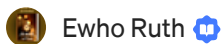



Oct 22



2



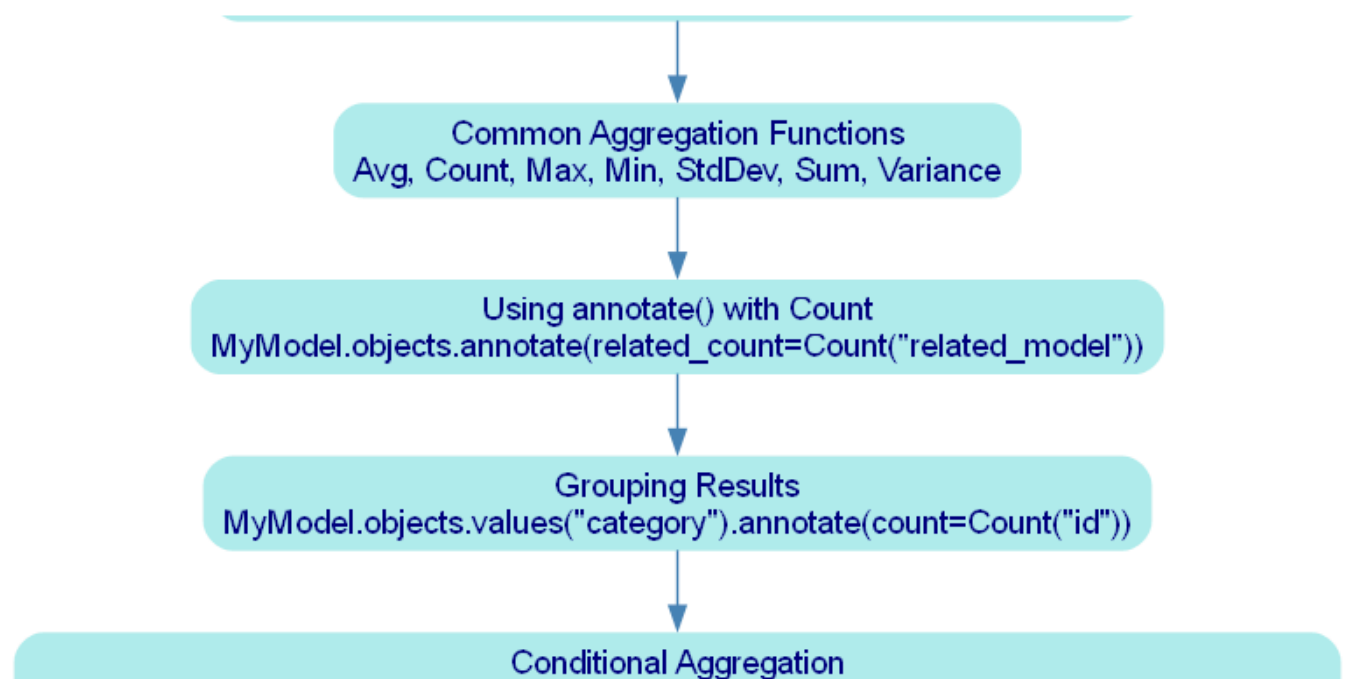



Ewho Ruth

“Mastering Python Decorators: A Comprehensive Guide for Enhancing Code Modularity and...”

1. Introduction

★ May 25, 2023 🖱 80



PY In Python in Plain English by Ewho Ruth


Django Aggregation Cheat Sheet: A Comprehensive Guide

Aggregation in Django

★ Oct 22 🖱 10

[See all from Ewho Ruth](#)

Recommended from Medium

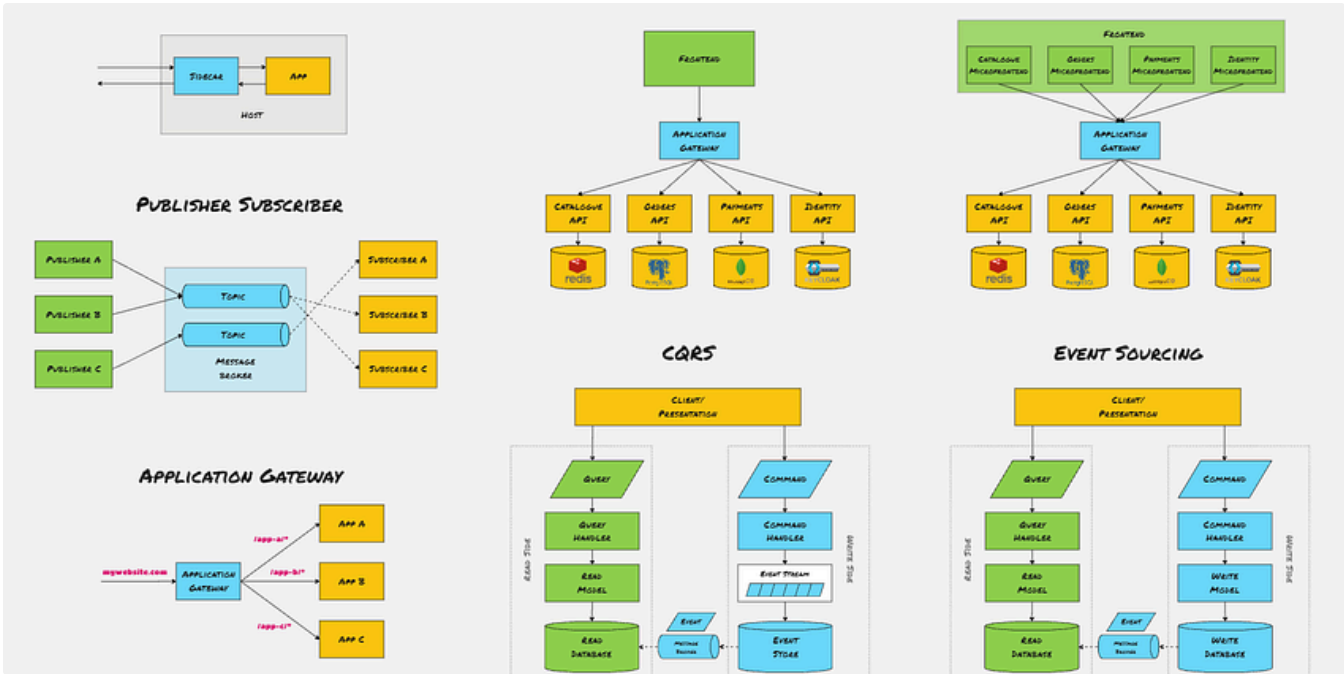
 In Stackademic by Abdur Rahman

Python is No More The King of Data Science

5 Reasons Why Python is Losing Its Crown

★ Oct 23 🖱 6.1K 💬 26





In Level Up Coding by Matt Bentley

My Favourite Software Architecture Patterns

Exploring my most loved Software Architecture patterns and their practical applications.

4d ago 2.6K 39



Lists




Coding & Development

11 stories · 902 saves



General Coding Knowledge

20 stories · 1740 saves



Predictive Modeling w/ Python

20 stories · 1675 saves



Stories to Help You Grow as a Software Developer

19 stories · 1471 saves



PY In Python in Plain English by KokaTic

Top Django Libraries for Enhancing Web Development Projects

Boost Your Django Projects with the Best Libraries Available

★ Oct 3 🖱️ 107 💬 2



django

REST

framework



Dhruv Ahuja

Stop Sending Messy Responses! The Ultimate Django DRF Response Structure You Need 🚀



Heads Up! Click here to unlock this article for free if you're not a Medium member!

★ Oct 30 🖱 67 💬 1



Harendra

How I Am Using a Lifetime 100% Free Server

Get a server with 24 GB RAM + 4 CPU + 200 GB Storage + Always Free

★ Oct 26 🖱 4.8K 💬 61



In Coding Beauty by Tari Ibaba

The new M4 Mac Mini is a MONSTER

How did Apple squeeze so much computing power into such a tiny space?

 6d ago  562  14



See more recommendations